

# End-to-End Training of Deep Visuomotor Policies

Sergey Levine\*, Chelsea Finn\*, Trevor Darrell, Pieter Abbeel.  
JMLR 17, 2016  
(335 cites)



# Visuomotor?


---

visuomotor 

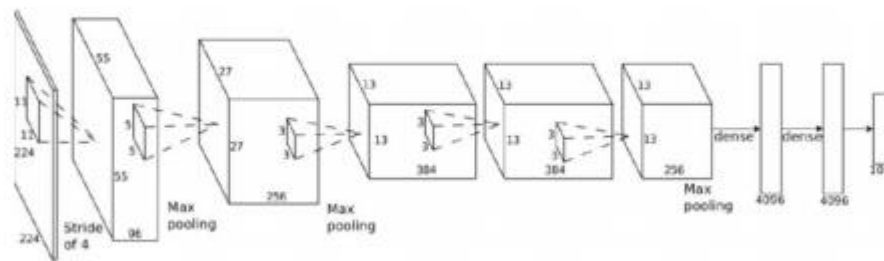
Also found in: [Dictionary](#).

vis·u·o·mo·tor (viz'yū-ō-mō'tōr),

Denoting the ability to synchronize visual information with physical movement, for example, driving a car or playing a video game of skill.

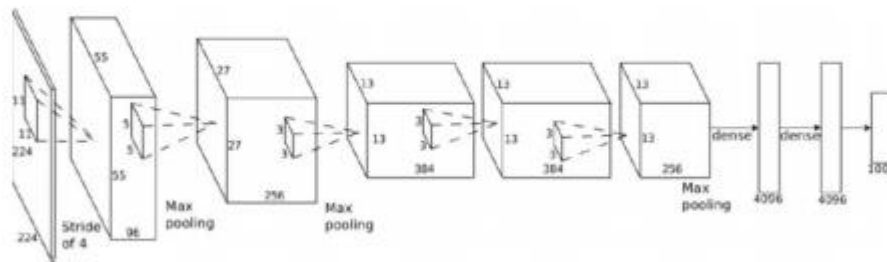
**"CITE"**  Farlex Partner Medical Dictionary © Farlex 2012

시각(카메라) 정보를 통해 신체(모터)를 움직이는 것이라고 봐도 무방할 듯



tiger	0.9
tiger cat	0.05
jaguar	0.02
lynx	0.01

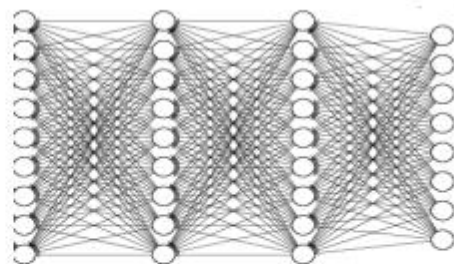
perception



tiger  
tiger  
tiger cat  
jaguar  
lynx

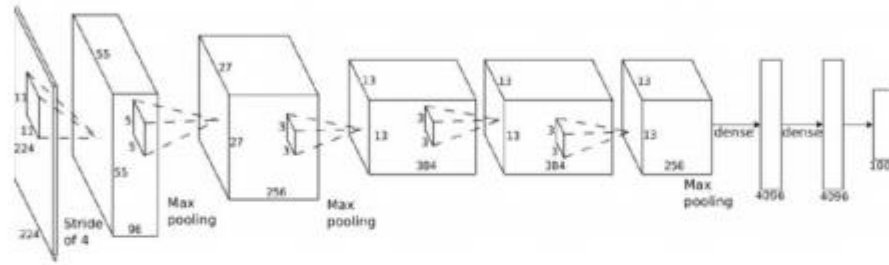


Action  
(run away)

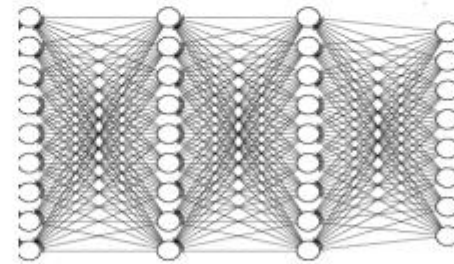


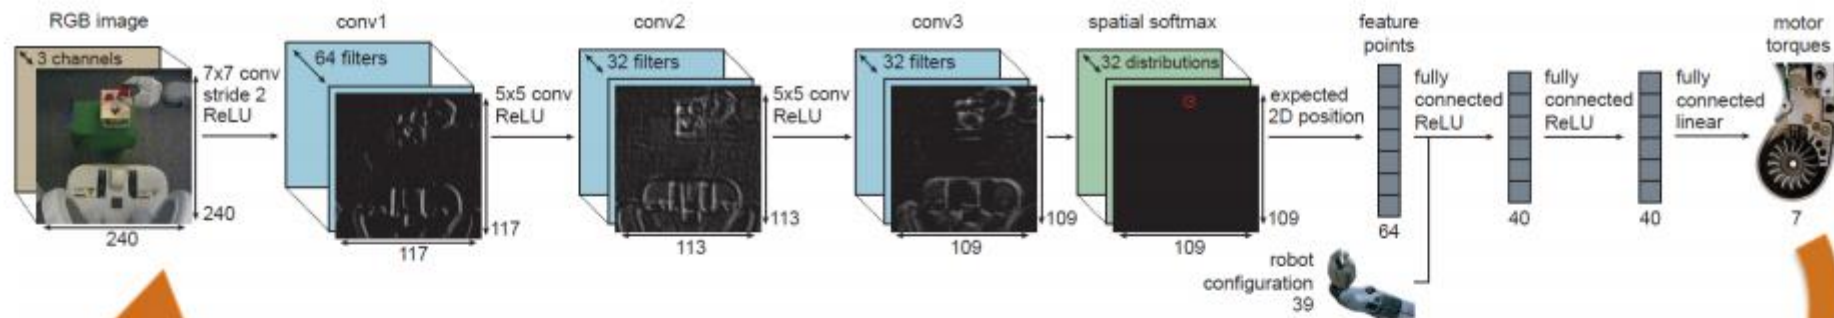
action

## sensorimotor loop



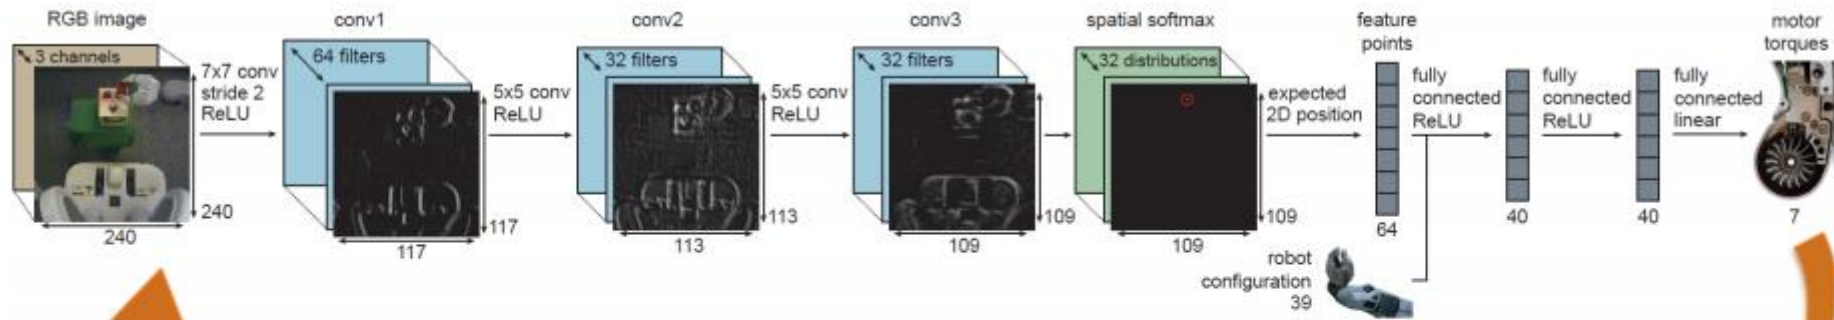
Action  
(run away)





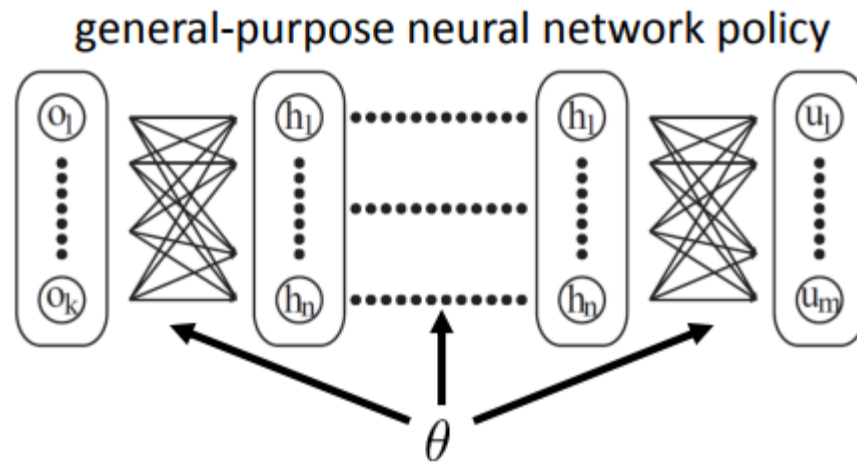
sensorimotor loop





sensorimotor loop

no direct supervision  
actions have consequences



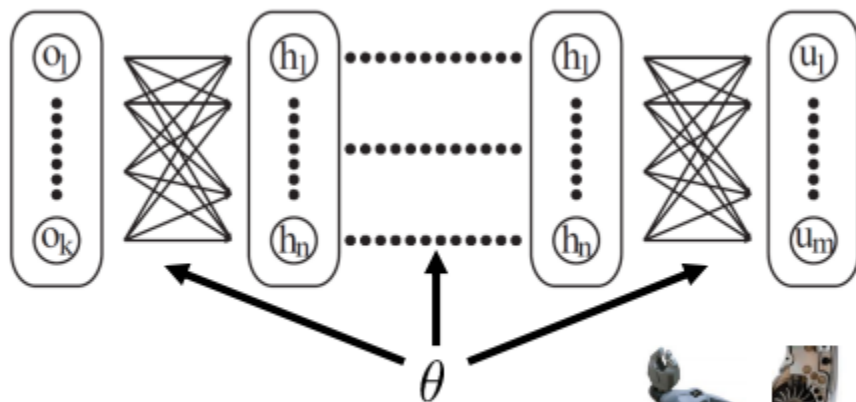
$$\theta = \arg \min_{\theta} E_{\pi_{\theta}} [\sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t)]$$

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$  – control policy

$\mathbf{o}_t$  – observation (may or may not be equal to  $\mathbf{x}_t$ )



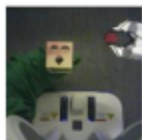
# general-purpose neural network policy



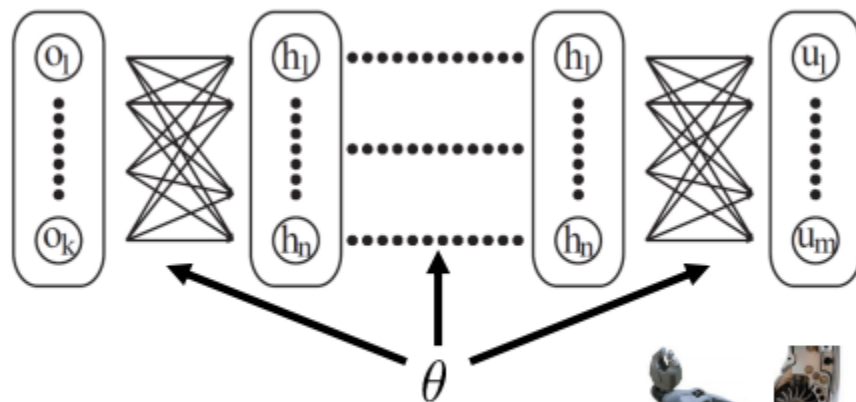
$$\theta = \arg \min_{\theta} E_{\pi_{\theta}} [\sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t)]$$

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$  – control policy

$\mathbf{o}_t$  – observation (may or may not be equal to  $\mathbf{x}_t$ )



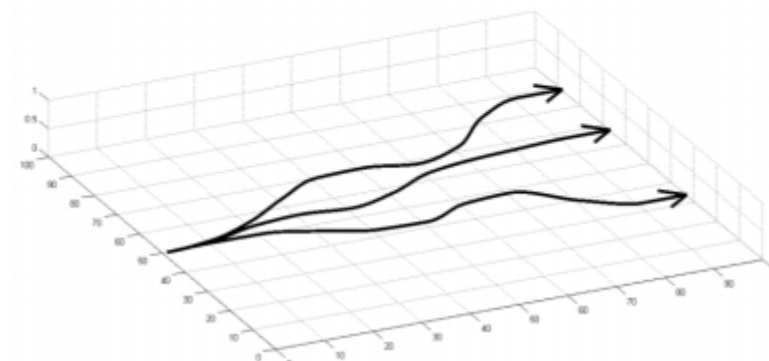
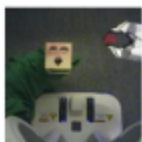
# general-purpose neural network policy



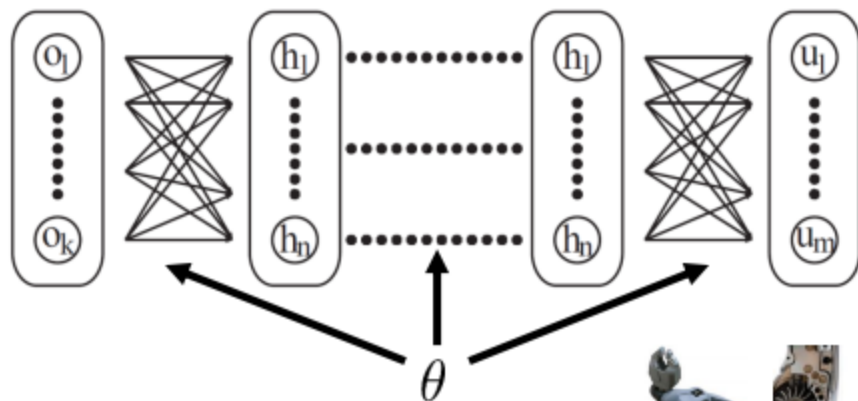
$$\theta = \arg \min_{\theta} E_{\pi_{\theta}} [\sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t)]$$

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$  – control policy

$\mathbf{o}_t$  – observation (may or may not be equal to  $\mathbf{x}_t$ )



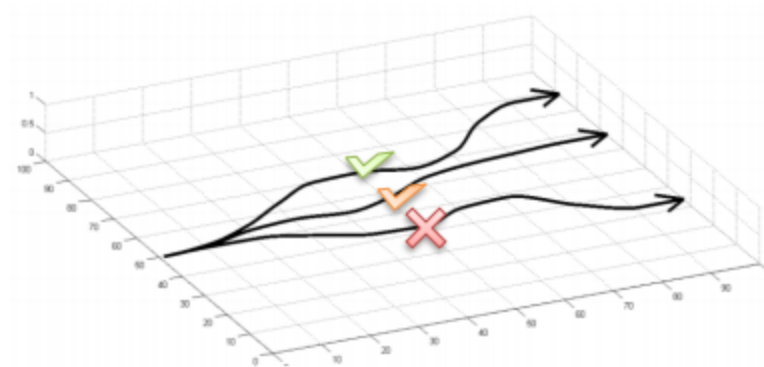
# general-purpose neural network policy



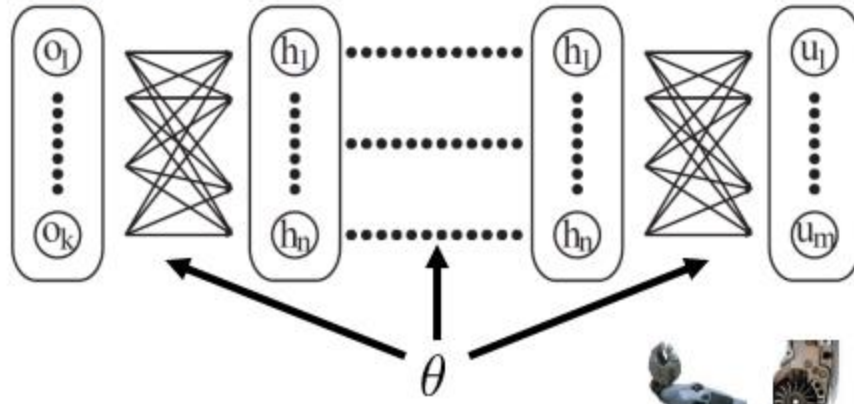
$$\theta = \arg \min_{\theta} E_{\pi_{\theta}} [\sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t)]$$

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$  – control policy

$\mathbf{o}_t$  – observation (may or may not be equal to  $\mathbf{x}_t$ )



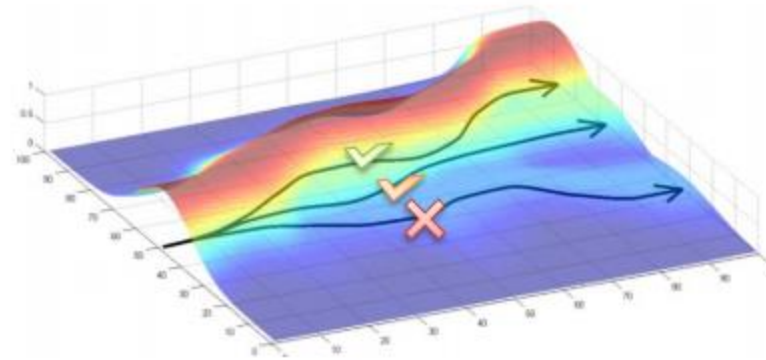
# general-purpose neural network policy



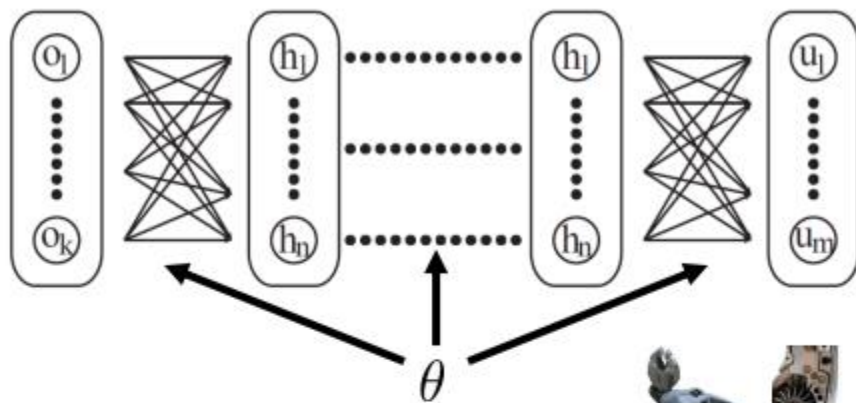
$$\theta = \arg \min_{\theta} E_{\pi_{\theta}} [\sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t)]$$

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$  – control policy

$\mathbf{o}_t$  – observation (may or may not be equal to  $\mathbf{x}_t$ )



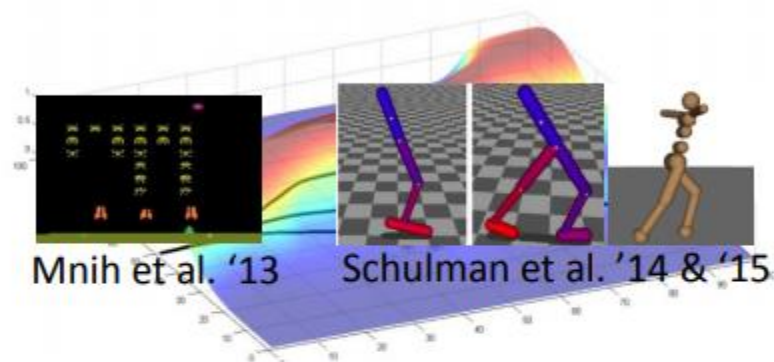
# general-purpose neural network policy



$$\theta = \arg \min_{\theta} E_{\pi_{\theta}} [\sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t)]$$

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$  – control policy

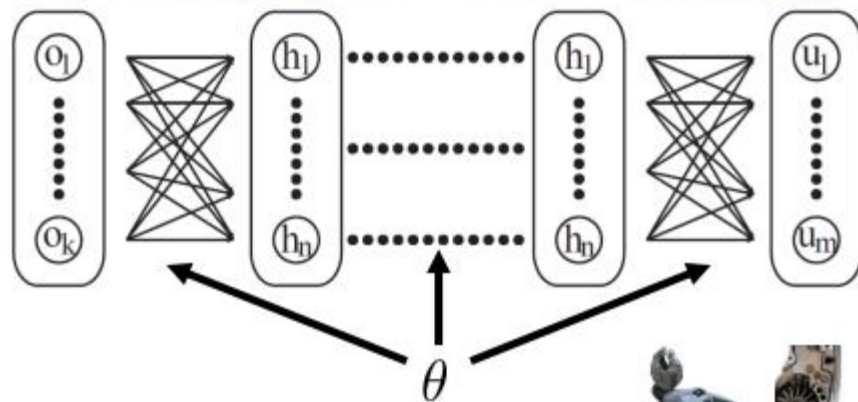
$\mathbf{o}_t$  – observation (may or may not be equal to  $\mathbf{x}_t$ )



Mnih et al. '13

Schulman et al. '14 & '15

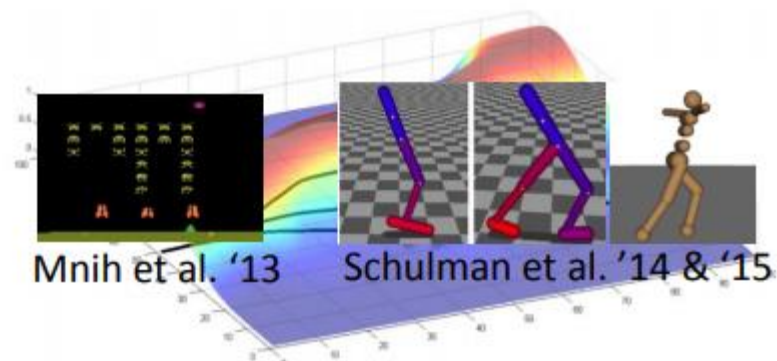
# general-purpose neural network policy



$$\theta = \arg \min_{\theta} E_{\pi_{\theta}} [\sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t)]$$

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$  – control policy

$\mathbf{o}_t$  – observation (may or may not be equal to  $\mathbf{x}_t$ )



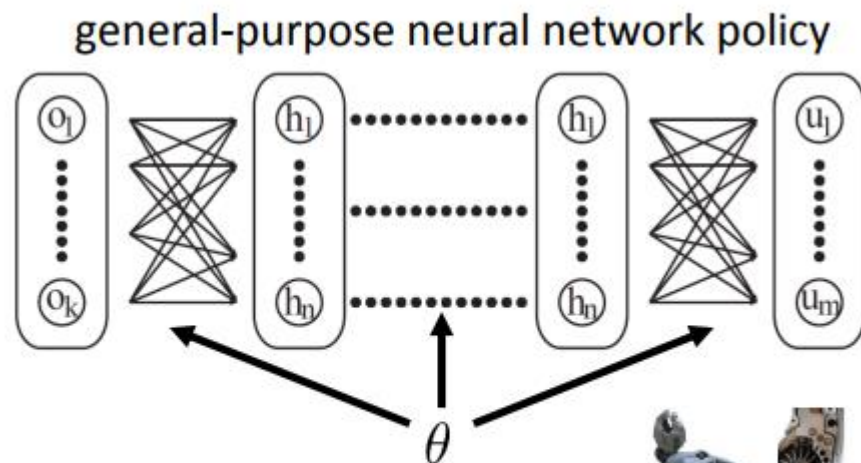
policy search (RL)

complex dynamics

complex policy

**HARD**

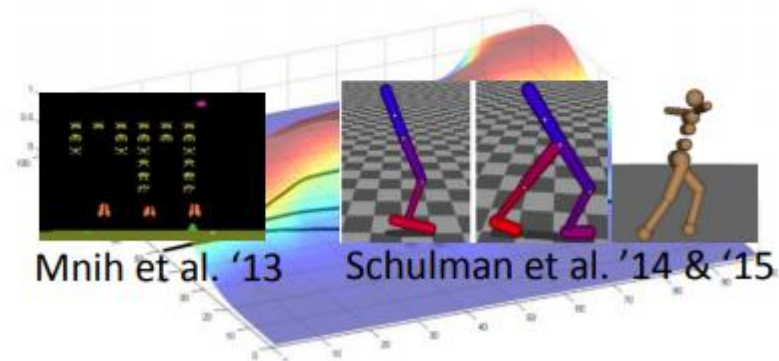




$$\theta = \arg \min_{\theta} E_{\pi_{\theta}} [\sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t)]$$

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$  – control policy

$\mathbf{o}_t$  – observation (may or may not be equal to  $\mathbf{x}_t$ )



policy search (RL)

complex dynamics

complex policy

**HARD**

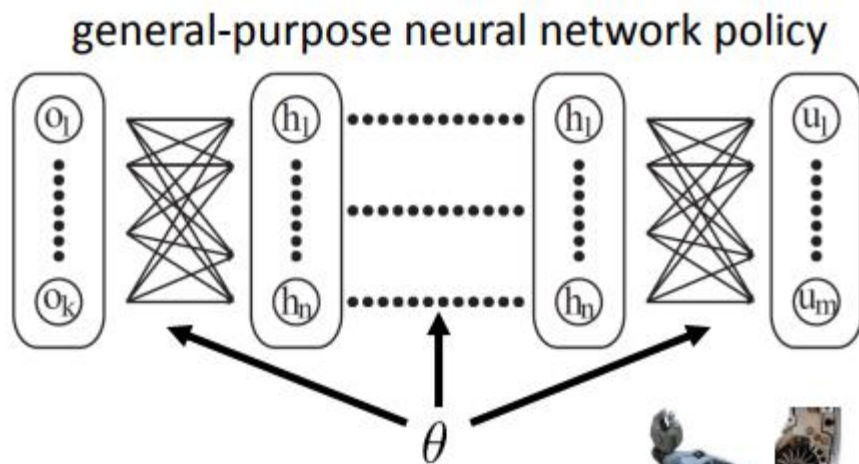
supervised learning

~~complex dynamics~~

complex policy

**EASY**

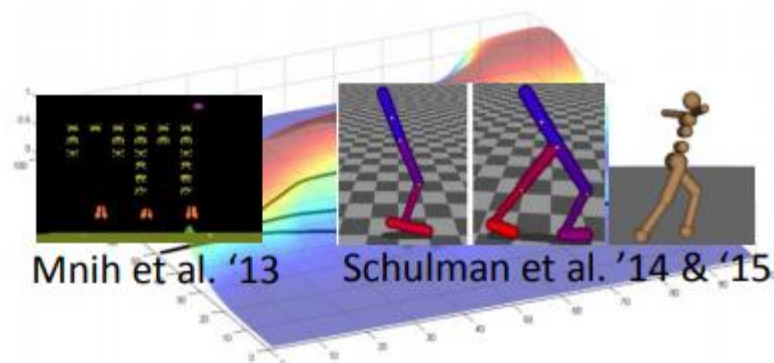




$$\theta = \arg \min_{\theta} E_{\pi_{\theta}} [\sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t)]$$

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$  – control policy

$\mathbf{o}_t$  – observation (may or may not be equal to  $\mathbf{x}_t$ )



policy search (RL)

complex dynamics

complex policy

**HARD**

supervised learning

~~complex dynamics~~

complex policy

**EASY**

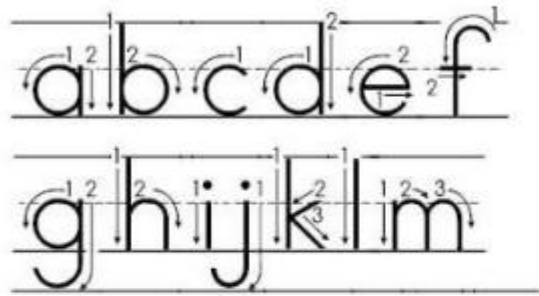
optimal control

complex dynamics

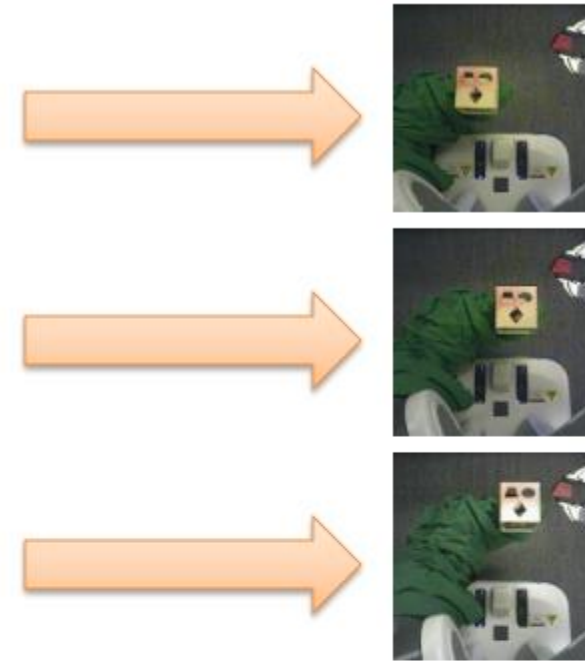
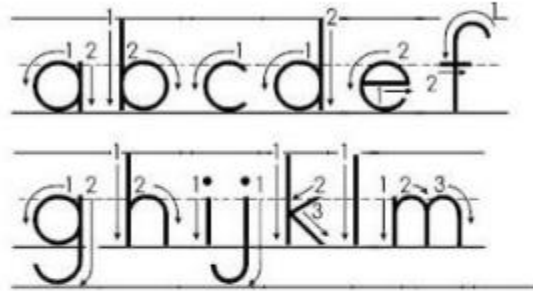
~~complex policy~~

**EASY**

1. break up the task:  
separately solve N  
different task instances

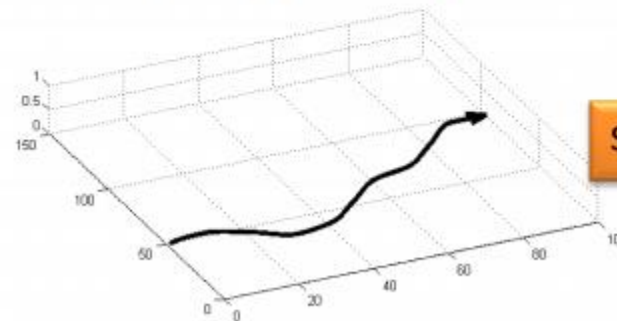


1. break up the task:  
separately solve N  
different task instances



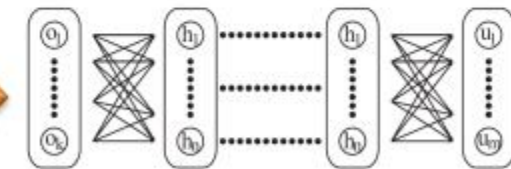
2. use supervised learning

trajectory-centric RL  
(fully observed)



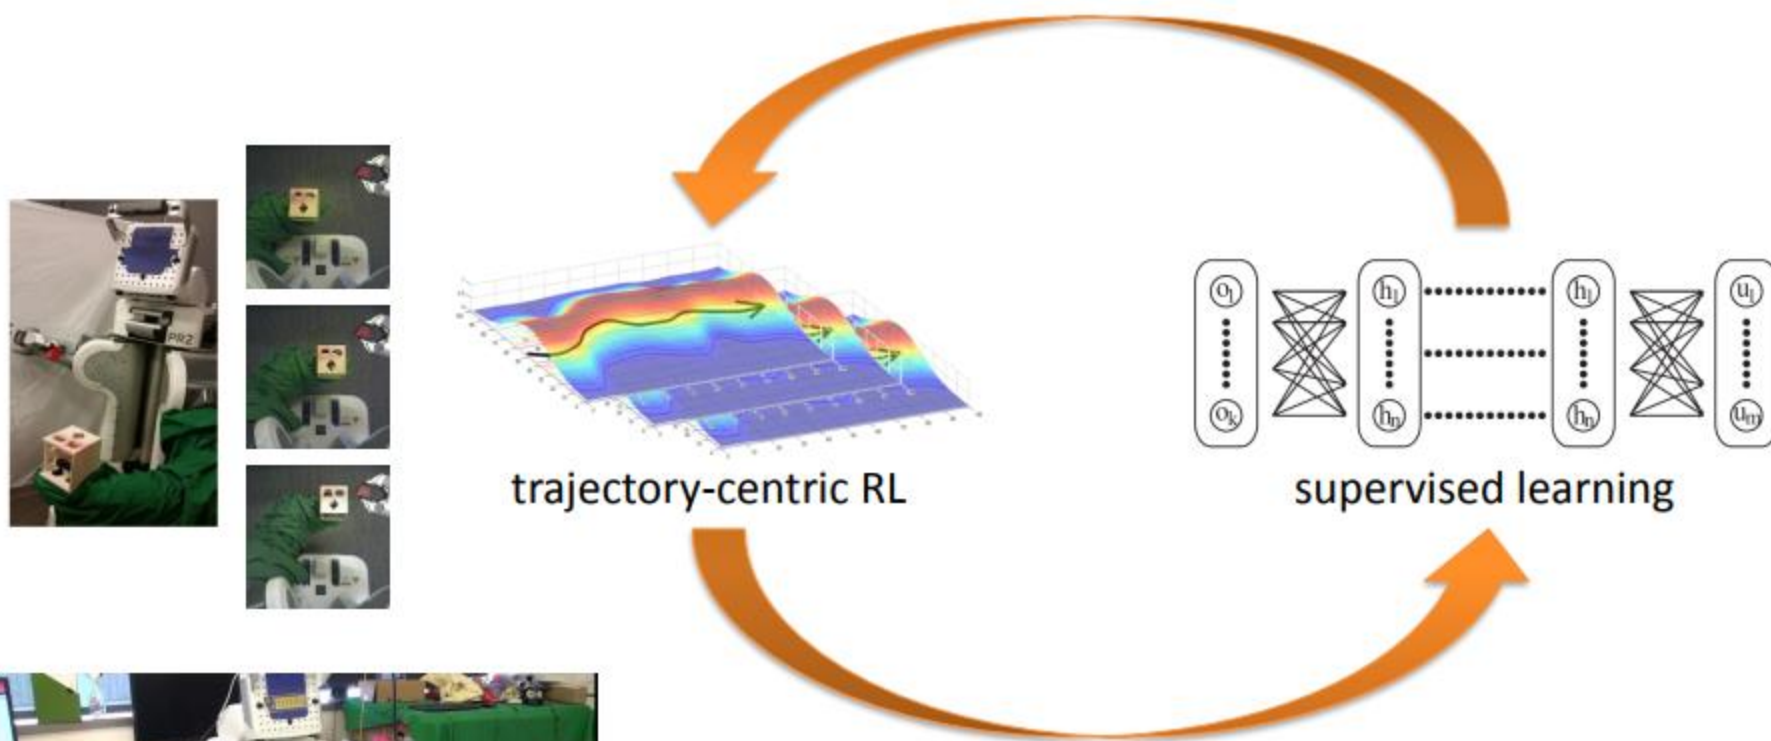
**state to action**

supervised learning



**observation to action**

# Guided Policy Search





expectation under  
current policy

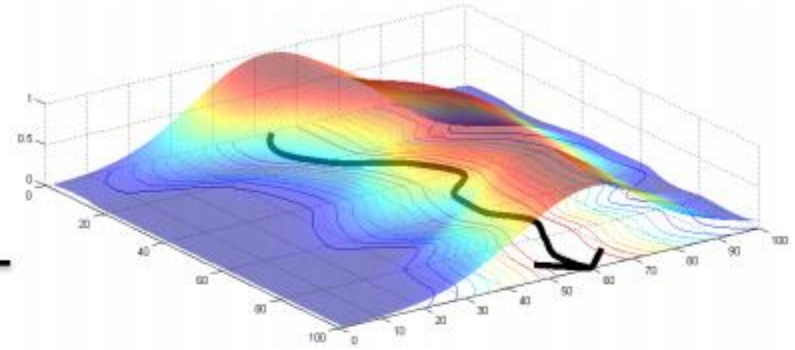
$$\min_{\theta} \overbrace{E_{\pi_{\theta}}[c(\tau)]}^{\text{expectation under current policy}}$$

↕

$$\min_{\theta, p(\tau)} E_p[c(\tau)] \quad \text{trajectory distribution(s)}$$


---


$$s.t. \pi_{\theta}(\mathbf{u}_t | \mathbf{o}(\mathbf{x}_t)) = p(\mathbf{u}_t | \mathbf{x}_t) \quad \forall t, \mathbf{x}_t, \mathbf{u}_t$$



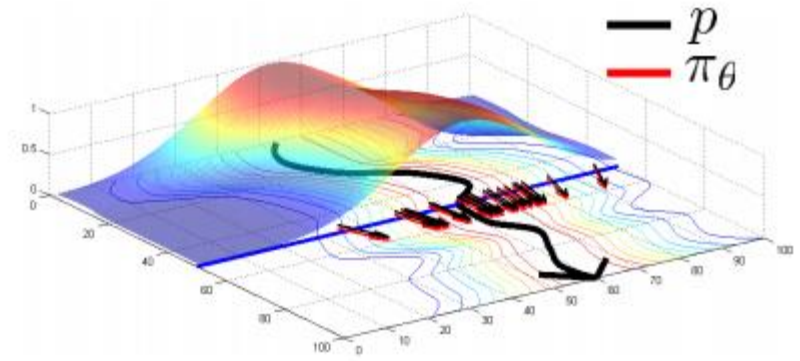
expectation under  
current policy

$$\min_{\theta} \overbrace{E_{\pi_{\theta}}[c(\tau)]}^{\text{expectation under current policy}}$$

↕

$$\min_{\theta, p(\tau)} E_p[c(\tau)]$$

$s.t. \pi_{\theta}(\mathbf{u}_t | \mathbf{o}(\mathbf{x}_t)) = p(\mathbf{u}_t | \mathbf{x}_t) \quad \forall t, \mathbf{x}_t, \mathbf{u}_t$



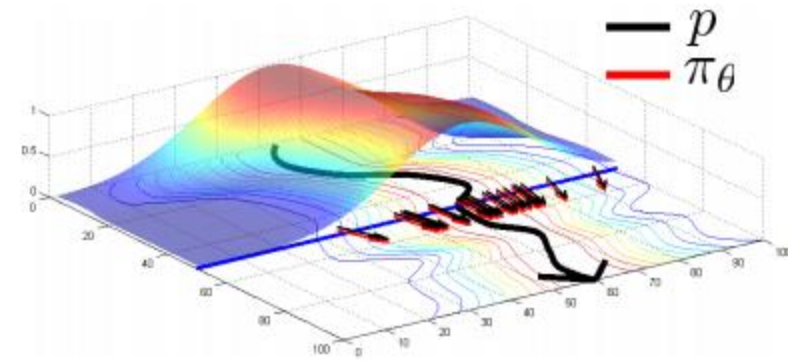
expectation under  
current policy

$$\min_{\theta} E_{\pi_{\theta}}[c(\tau)]$$

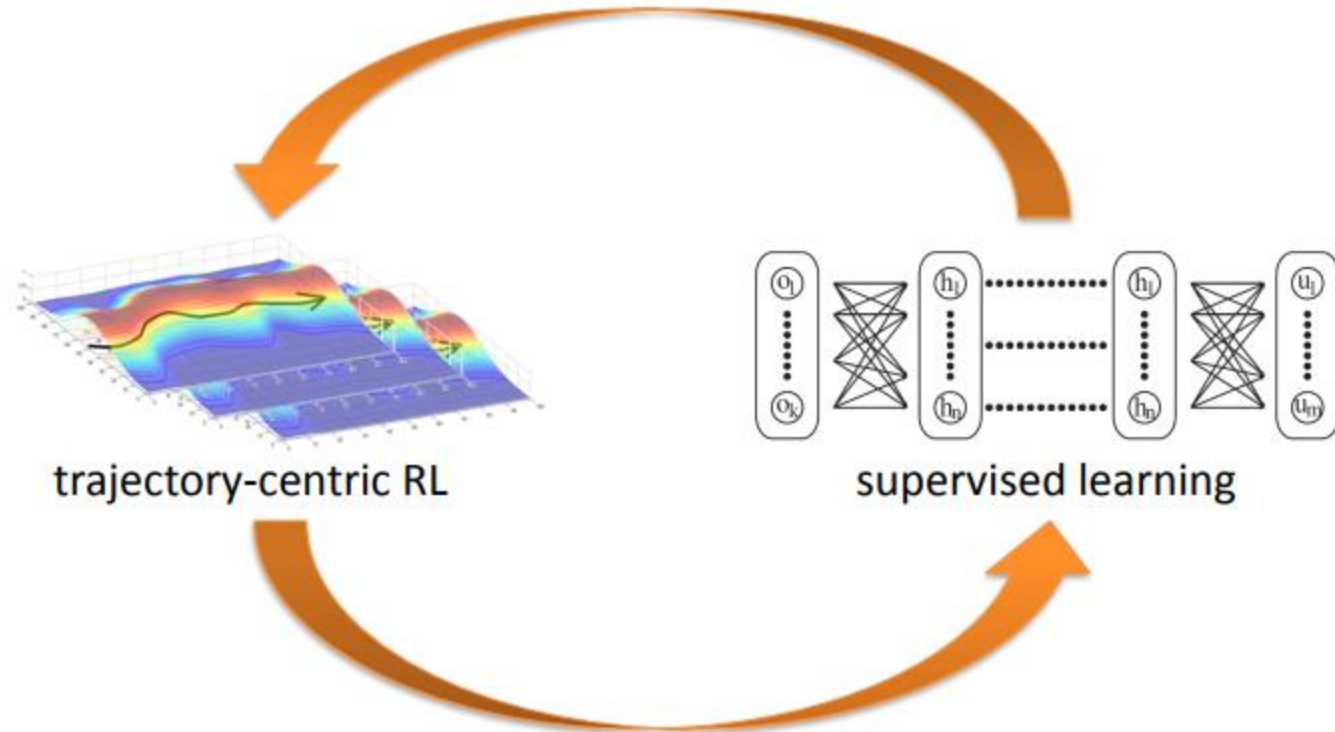


$$\min_{\theta, p(\tau)} E_p[c(\tau)]$$

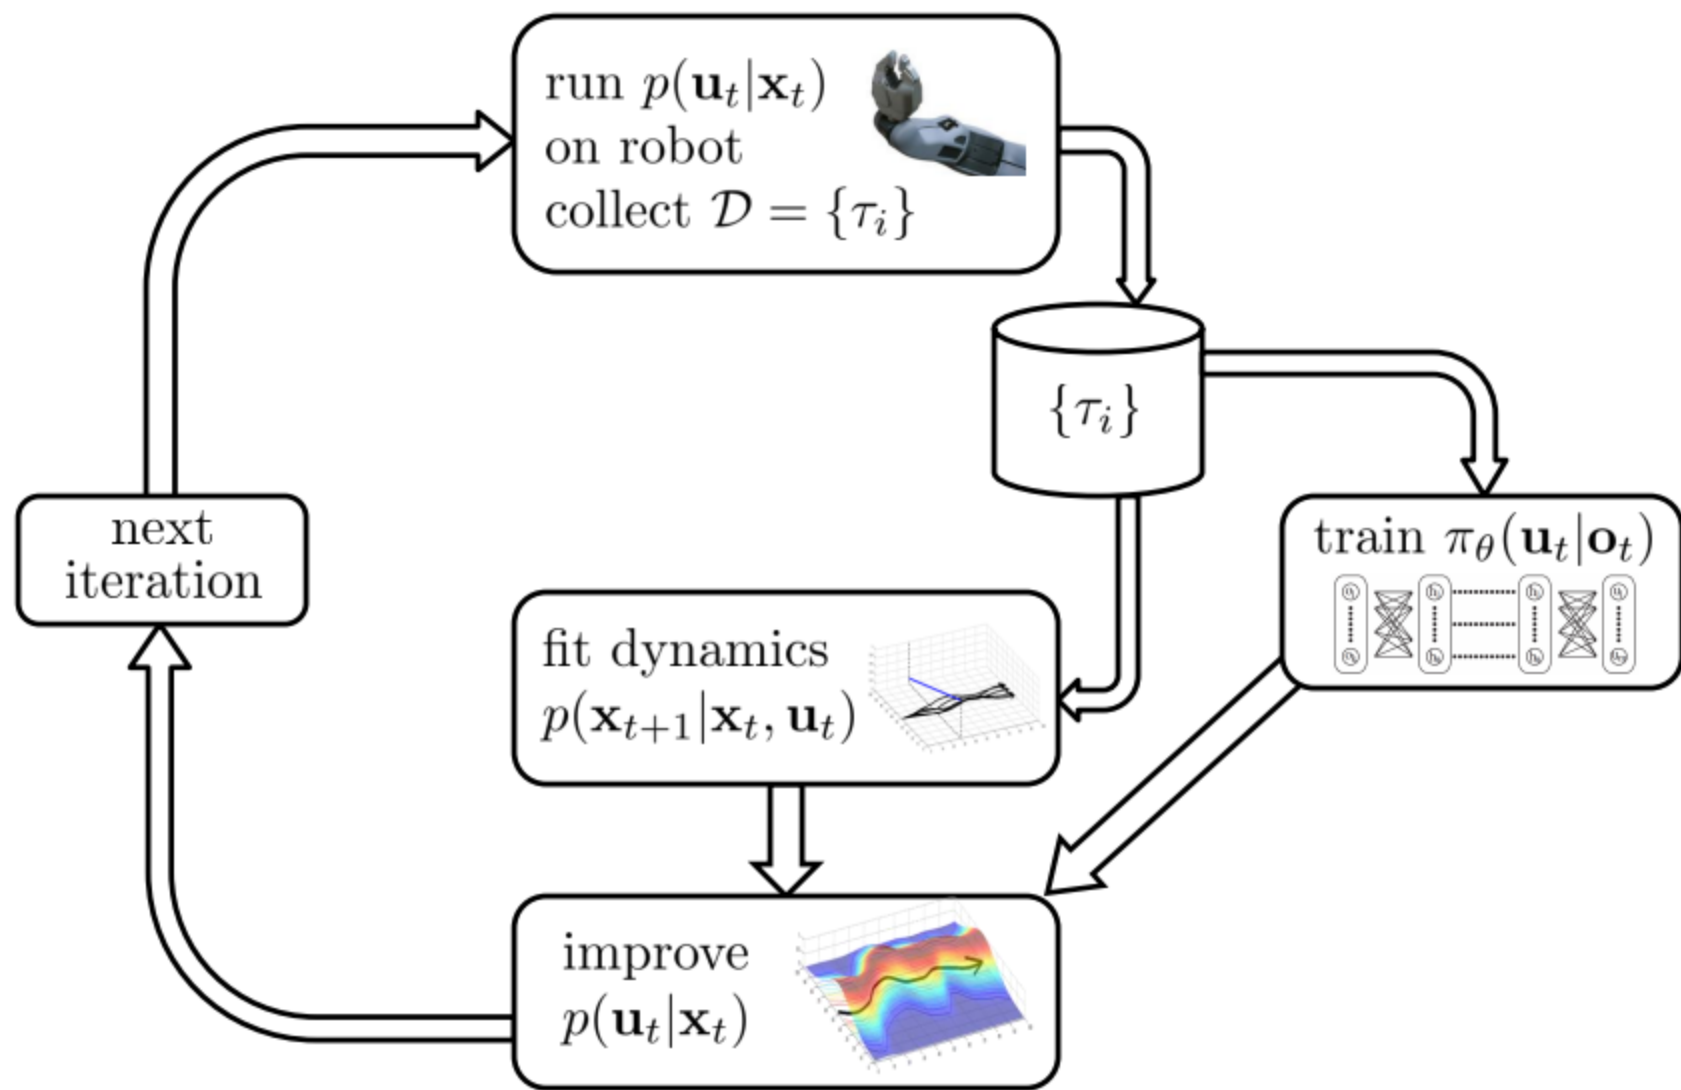
$$s.t. \pi_{\theta}(\mathbf{u}_t | \mathbf{o}(\mathbf{x}_t)) = p(\mathbf{u}_t | \mathbf{x}_t) \quad \forall t, \mathbf{x}_t, \mathbf{u}_t$$

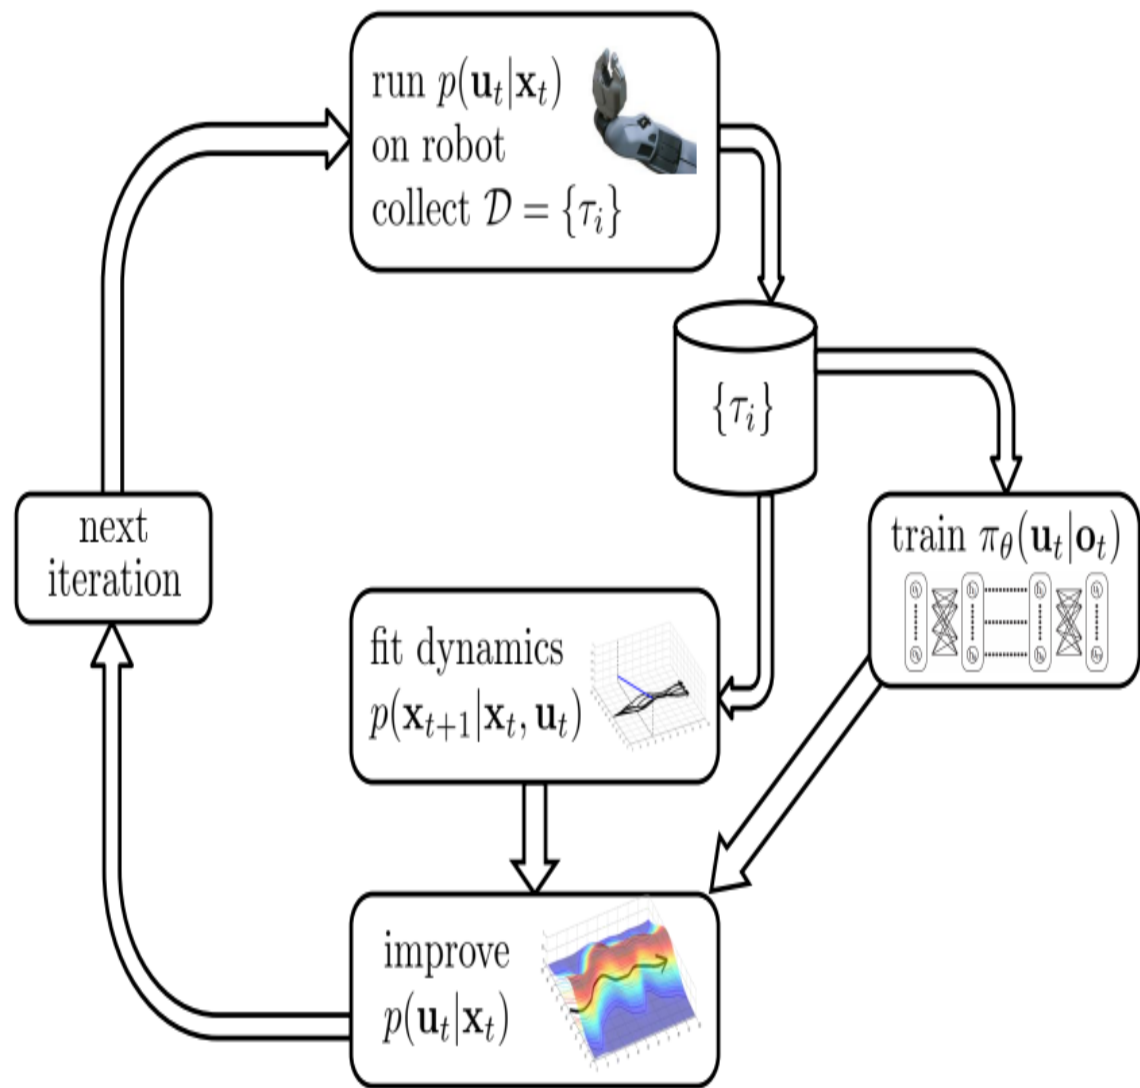


solve using Bregman ADMM (BADMM), a type of dual decomposition method



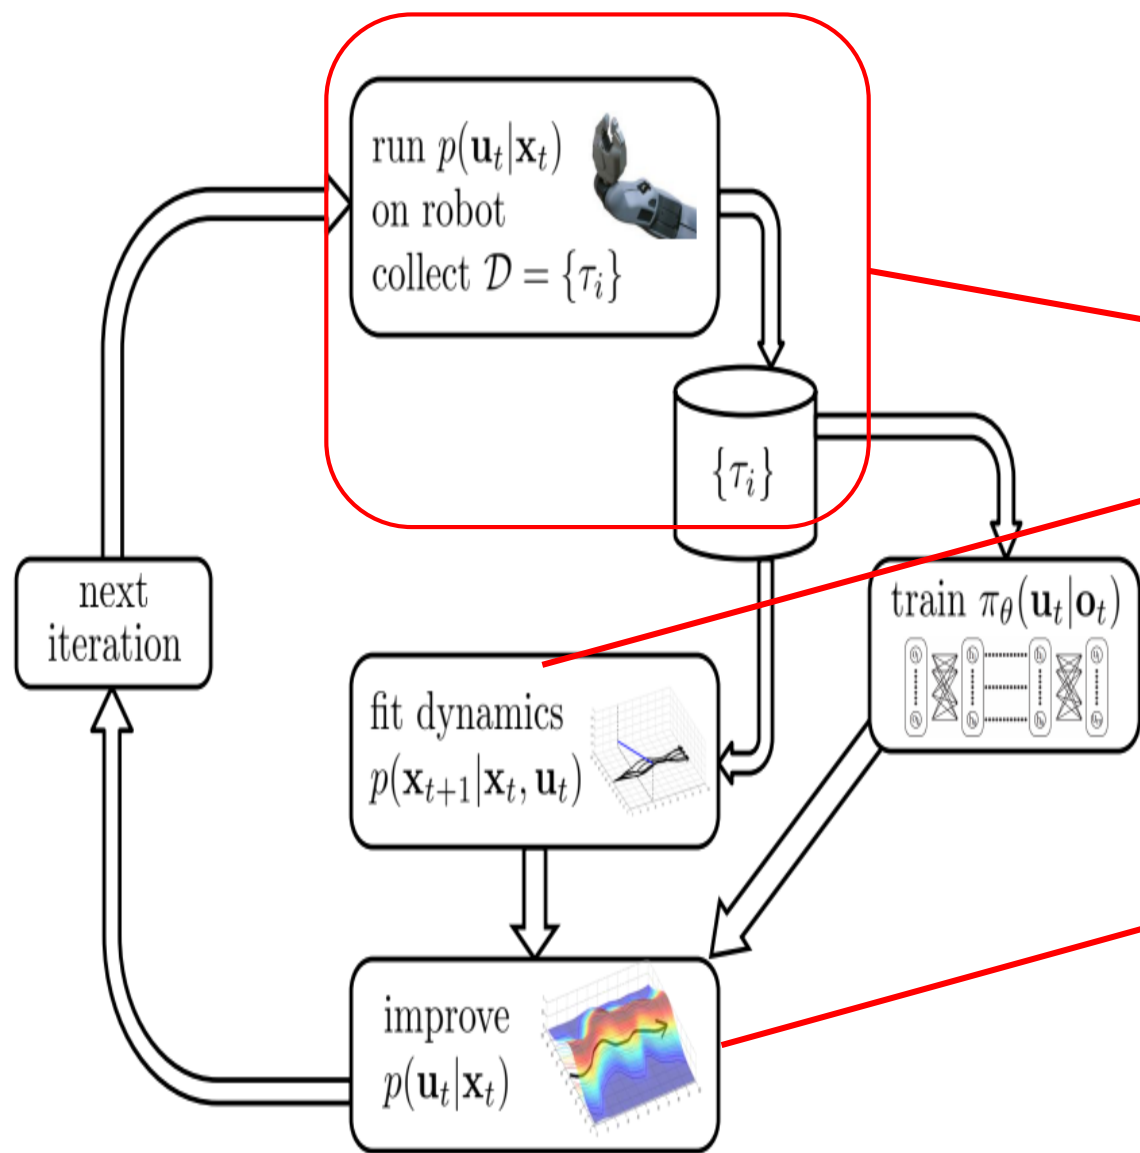






## Algorithm 1 Guided Policy Search

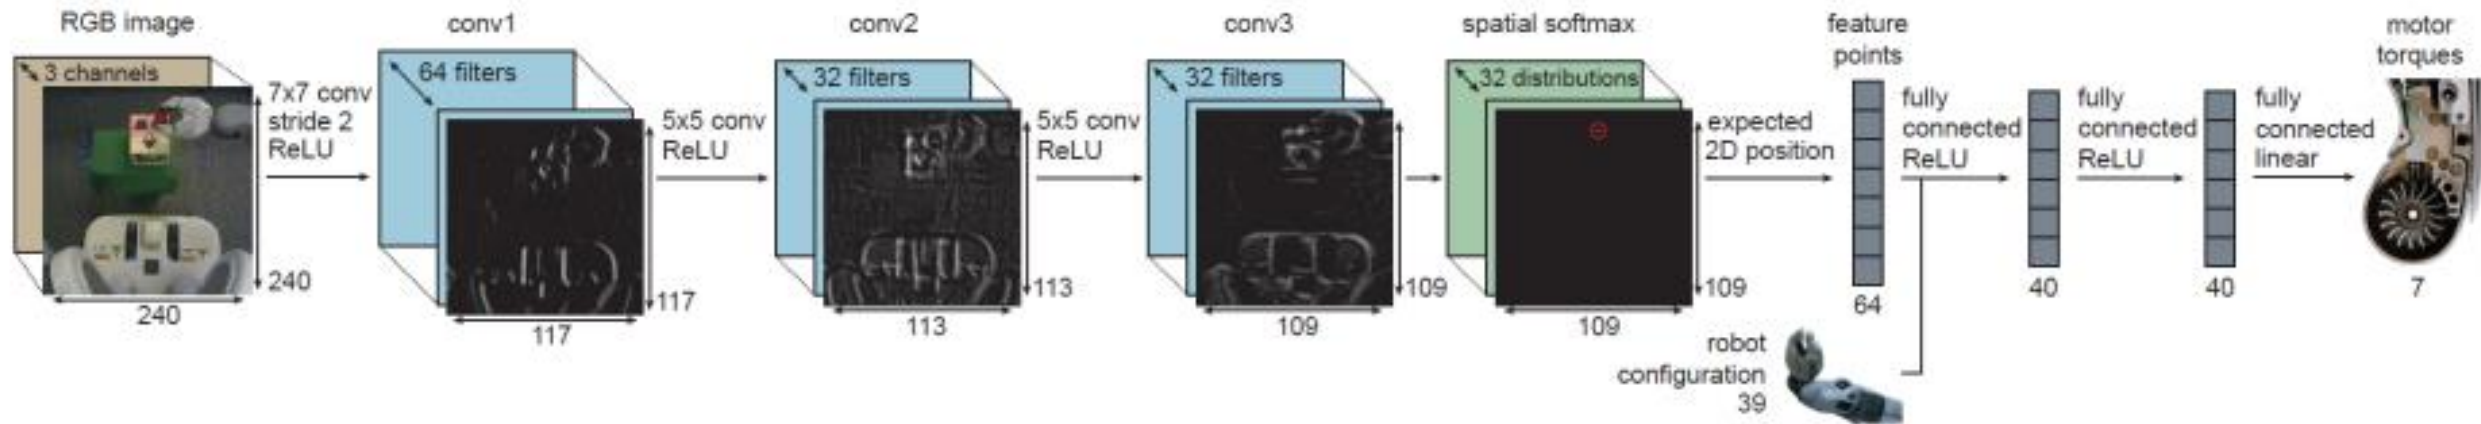
- 1: Generate DDP solutions  $\pi_{\mathcal{G}_1}, \dots, \pi_{\mathcal{G}_n}$
- 2: Sample  $\zeta_1, \dots, \zeta_m$  from  $q(\zeta) = \frac{1}{n} \sum_i \pi_{\mathcal{G}_i}(\zeta)$
- 3: Initialize  $\theta^* \leftarrow \arg \max_{\theta} \sum_i \log \pi_{\theta^*}(\zeta_i)$
- 4: Build initial sample set  $\mathcal{S}$  from  $\pi_{\mathcal{G}_1}, \dots, \pi_{\mathcal{G}_n}, \pi_{\theta^*}$
- 5: **for** iteration  $k = 1$  to  $K$  **do**
- 6:   Choose current sample set  $\mathcal{S}_k \subset \mathcal{S}$
- 7:   Optimize  $\theta_k \leftarrow \arg \max_{\theta} \Phi_{\mathcal{S}_k}(\theta)$
- 8:   Append samples from  $\pi_{\theta_k}$  to  $\mathcal{S}_k$  and  $\mathcal{S}$
- 9:   Optionally generate adaptive guiding samples
- 10:   Estimate the values of  $\pi_{\theta_k}$  and  $\pi_{\theta^*}$  using  $\mathcal{S}_k$
- 11:   **if**  $\pi_{\theta_k}$  is better than  $\pi_{\theta^*}$  **then**
- 12:     Set  $\theta^* \leftarrow \theta_k$
- 13:     Decrease  $w_r$
- 14:   **else**
- 15:     Increase  $w_r$
- 16:     Optionally, resample from  $\pi_{\theta^*}$
- 17:   **end if**
- 18: **end for**
- 19: Return the best policy  $\pi_{\theta^*}$

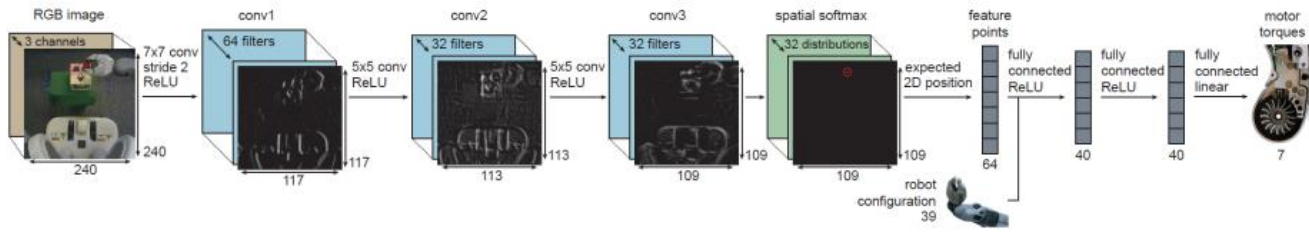


## Algorithm 1 Guided Policy Search

- 1: Generate DDP solutions  $\pi_{\mathcal{G}_1}, \dots, \pi_{\mathcal{G}_n}$
- 2: Sample  $\zeta_1, \dots, \zeta_m$  from  $q(\zeta) = \frac{1}{n} \sum_i \pi_{\mathcal{G}_i}(\zeta)$
- 3: Initialize  $\theta^* \leftarrow \arg \max_{\theta} \sum_i \log \pi_{\theta^*}(\zeta_i)$
- 4: Build initial sample set  $\mathcal{S}$  from  $\pi_{\mathcal{G}_1}, \dots, \pi_{\mathcal{G}_n}, \pi_{\theta^*}$
- 5: **for** iteration  $k = 1$  to  $K$  **do**
- 6:   Choose current sample set  $\mathcal{S}_k \subset \mathcal{S}$
- 7:   Optimize  $\theta_k \leftarrow \arg \max_{\theta} \Phi_{\mathcal{S}_k}(\theta)$
- 8:   Append samples from  $\pi_{\theta_k}$  to  $\mathcal{S}_k$  and  $\mathcal{S}$
- 9:   Optionally generate adaptive guiding samples
- 10:   Estimate the values of  $\pi_{\theta_k}$  and  $\pi_{\theta^*}$  using  $\mathcal{S}_k$
- 11:   **if**  $\pi_{\theta_k}$  is better than  $\pi_{\theta^*}$  **then**
- 12:     Set  $\theta^* \leftarrow \theta_k$
- 13:     Decrease  $w_r$
- 14:   **else**
- 15:     Increase  $w_r$
- 16:     Optionally, resample from  $\pi_{\theta^*}$
- 17:   **end if**
- 18: **end for**
- 19: Return the best policy  $\pi_{\theta^*}$

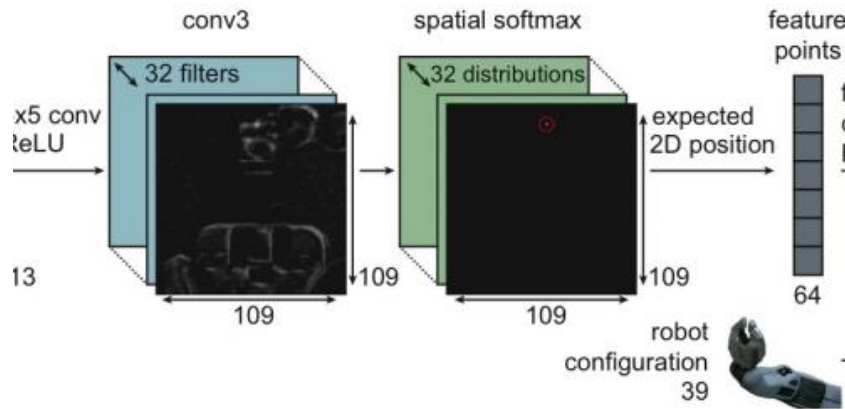
# Network architecture





spatial softmax 라고 읽고 spatial soft-argmax라 부른다

1. 채널별로 softmax를 구함
2. 채널별로 구한 softmax 결과와 x, y 좌표 값을 곱한 후 더한다.

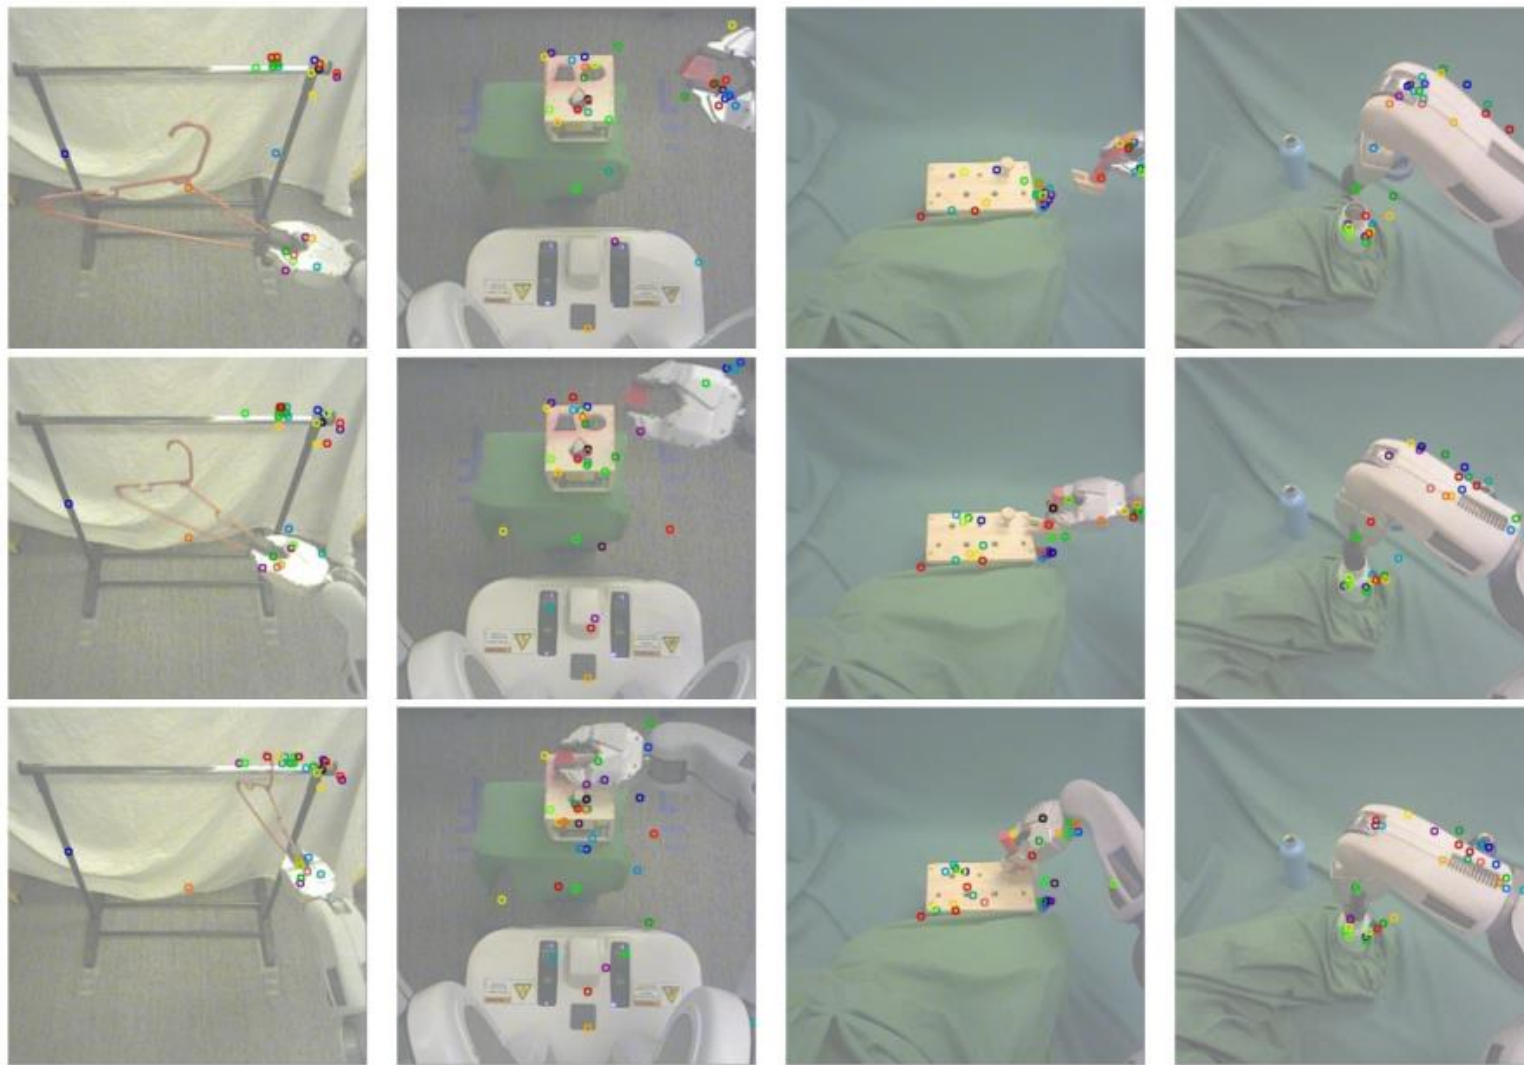


$$s_{cij} = e^{a_{cij}} / \sum_{i'j'} e^{a_{ci'j'}}$$

$$f_{cx} = \sum_{ij} s_{cij} x_{ij} \text{ and } f_{cy} = \sum_{ij} s_{cij} y_{ij}$$

$$\text{softargmax}(x) = \sum_i \frac{e^{\beta x_i}}{\sum_j e^{\beta x_j}} i$$





(a) hanger

(b) cube

(c) hammer

(d) bottle

Figure 10: Feature points tracked by the policy during task execution for each of the four tasks. Each feature point is displayed in a different random color, with consistent coloring across images. The policy finds features on the target object and the robot gripper and arm. In the bottle cap task, note that the policy correctly ignores the distractor bottle in the background, even though it was not present during training.



**Position 2**

real time

autonomous execution



# Training

- SL + RL
- Cost function